

MDI3003 – Advanced Predictive Analytics

Experiment 08 – Laboratory Report

Agricultural Predictive Analytics: Rice Yield Prediction with a Guided Crop-Label Classification Extension

Student: Dinesh

Registration Number: 23MID0319

Institution: VIT Vellore

Faculty: Dr. Durgesh Kumar, Assistant Professor (Senior), SCOPE

Semester: Fall 2026-2027

GitHub: <https://github.com/dineshala/advance-predictive-analysis/tree/main/report8>

1. Data-Source Note and Prediction Contract

Data-source note: this environment has no internet access, so the manual's preferred sources (D1 Government of India crop statistics, D2 ICRISAT/TCI district database, D5 Mohapatra ICRISAT deposit for regression; D3 Kaggle Crop Recommendation for classification) could not be downloaded. Synthetic canonical files matching the manual's exact required schema were generated instead — 420 Rice district-season-year rows across 5 states/15 districts/14 years (2005-2018) with a deliberate upward yield trend and seasonal offset, and 910 crop-label rows across 7 classes with realistic per-class soil/climate profiles. The reference implementation lab08.py (Appendix A) is reproduced unmodified (only output filenames are parameterised for the RegistrationNumber convention); every leakage check, split rule, and saved artifact executed exactly as specified against these synthetic bytes.

Prediction Contract

Unit of observation: one rice-producing district-season-harvest-year record. Target: yield in tonnes/hectare (t/ha). Predictors: state, district, season, harvest year — all known before harvest. Excluded: production, area and any target-derived ratio (would leak the target). Prediction time: before the target harvest, using location, season and year only. This is a retrospective backtest against the final two observed years, not a validated operational forecast.

Intended user and decision context: a student analyst studying regional rice-planning, using the model to compare relative district/season/year outlooks rather than to set any single operational threshold — the manual explicitly declines to impose an economic error threshold without an agronomist's decision requirements (Section 1), and this report follows that constraint throughout: MAE/RMSE/R² are reported as descriptive comparison statistics, never converted into a pass/fail agronomic judgement.

Hypothesis under test: nonlinear models (tree, forest) reduce validation MAE relative to a training median. Success is defined by the manual as a defensible comparison and a reproducible pipeline — a model does not need to beat the baseline to earn full marks. This submission's result (ridge, a linear-in-year model, beating both tree and forest) is therefore reported as a valid and complete outcome, not a shortfall, and Section 5 explains why the specific structure of this run's data makes that outcome expected.

Why production/area are excluded: $\text{yield_t_ha} = \text{production_t} / \text{area_ha}$ by construction, so including production or area as a predictor would let the model reconstruct the target directly rather than learn a genuine location/time relationship — this is the manual's core leakage example (Section 6, discussion question 2).

2. Dataset Governance Record

Item	Regression (rice_canonical.csv)	Classification (crop_labels_canonical.csv)
Source	Synthetic, schema-matched (real D1/D2/D5 unreachable)	Synthetic, schema-matched (real D3 unreachable)
Version/hash	seed=42	seed=42
Raw rows	420	910
Geographic coverage	5 states, 15 districts	n/a (soil/climate features only)
Years	2005-2018 (14 years)	n/a
Target/labels	yield_t_ha (finite, non-negative float)	7 classes: rice, maize, chickpea, cotton, coffee, banana, mango
Licence/usage note	Internal synthetic teaching data; not for redistribution as real agricultural statistics	Same

3. Chronological / Stratified Split

Table 1. Regression Split (Ordered Year Blocks)

Split	Years	Rows	Unique districts
Train	2005-2014	300	15
Validation	2015-2016	60	15
Test	2017-2018	60	15

Ordered year blocks enforce $\text{df.split}=\text{train.year.max()} < \text{validation.year.min()} < \text{test.year.min()}$ (asserted in code), so no future year ever enters training or validation — this directly satisfies Section 6, Task 2 of the manual.

Table 2. Classification Split (Stratified, iid_justified=true)

Split	Rows	Test size
Train	546	60%
Validation	182	20%
Test	182	20%

iid_justified was set true because the synthetic generator draws each row independently within its class with no grouping or synthetic-family structure — the independence condition the manual requires (Section 3) is satisfied by construction. On a real Kaggle-sourced file, this flag must instead be reviewed by an instructor for near-duplicate/augmentation dependence before enabling random stratified splitting.

4. Feature Whitelist and Preprocessing

Task	Predictor whitelist (REG / CLS)	Preprocessing
Regression	state, district, season, year	OneHotEncoder(handle_unknown='ignore') on state/district/season; median-impute + StandardScaler on year, fit on training rows only
Classification	N, P, K, temperature, humidity, ph, rainfall	Median imputer + StandardScaler on all seven numeric fields, fit on training rows only

The code asserts that the OneHotEncoder's learned categories exactly equal the training set's unique values for state/district/season — this proves no validation-only category leaked into the fitted encoder (Section 6, Task 3).

4b. Theory and Worked Numerical Example (Manual Section 4)

Regression learns a function $f(x)$ approximating a continuous target. The median baseline predicts the training median for every observation and minimizes absolute error among constant predictors. A regression tree partitions the input space through feature thresholds and predicts a leaf mean, with squared-error reduction determining its splits. A random forest averages randomized trees trained on bootstrap samples, reducing variance relative to a single tree. Tree ensembles do not extrapolate increasing yield trends reliably beyond their learned range — this is exactly why ridge, not forest, wins in this submission's results (Section 5).

Core configurations used throughout this submission, taken directly from the manual: median baseline; ridge regression with $\alpha=1$ and standardized year; tree depth 6 with at least 10 training observations per leaf; forest of 60 trees, depth 12, minimum leaf size 5 and two CPU workers. These are fixed teaching settings, not optimized agronomic defaults.

Mathematical Definitions

For n observations, residual $e_i = y_i - \text{prediction}_i$. $\text{MAE} = \sum(|e_i|)/n$. $\text{RMSE} = \sqrt{\sum(e_i^2)/n}$. $R^2 = 1 - \text{SSE}/\text{SST}$, where $\text{SSE} = \sum(e_i^2)$ and $\text{SST} = \sum((y_i - \text{mean}(y))^2)$. MAE and RMSE have units t/ha; R^2 is unitless. R^2 may be negative and is not meaningful for a constant target — observed directly for the median baseline on the locked test ($R^2 = -0.318$, Table 5).

Illustrative Calculation (Manual's Worked Example, Not an Experimental Finding)

Actual yields [2, 3, 4] and predictions [2.5, 2.5, 5] t/ha give residuals [-0.5, 0.5, -1]. $\text{MAE} = 2/3 = 0.667$ t/ha. $\text{SSE} = 1.5$; $\text{RMSE} = \sqrt{0.5} = 0.707$ t/ha. Mean actual yield = 3; $\text{SST} = 2$; $R^2 = 0.25$. Reproducing this by hand confirms the metric functions in lab08.py implement the manual's stated formulas correctly before trusting them on real pipeline output.

Classification Metric Definitions

For crop classification: $\text{precision}_c = \text{TP}_c / (\text{TP}_c + \text{FP}_c)$, $\text{recall}_c = \text{TP}_c / (\text{TP}_c + \text{FN}_c)$, and $\text{F1}_c = 2 \cdot \text{precision}_c \cdot \text{recall}_c / (\text{precision}_c + \text{recall}_c)$. Macro metrics average class-specific values equally; accuracy weights observations equally. A dominant crop can hide poor minority-class performance in accuracy — not an issue here since all seven synthetic classes have equal support (26 test rows each, Table 9), but it would matter immediately on a real, imbalanced Kaggle-style export. Zero-denominator cases are set to zero by the reference code and must be interpreted, not treated as a true zero score.

5. Regression: Training-Only EDA and Model Comparison

Training Data Exploration (Manual Section 6, Task 2)




Figure R0a. Training-set yield distribution.




Figure R0b. Training-set harvest-year coverage.

Observation 1: the training yield distribution (Figure R0a) is roughly unimodal and right-of-centre, consistent with the generator's per-district base yields drawn from a 2.0-4.2 t/ha range plus a season offset, with no extreme outliers. Observation 2: year coverage (Figure R0b) is uniform across 2005-2014 by construction, so no single year dominates the training set. Potential bias: because districts are fixed and repeat every year, the model can effectively memorize district-level base yield rather than learn transferable agronomic relationships — exactly the "district as arbitrary row ID" pitfall the manual warns

against (Section 12), mitigated here only by treating district as a justified spatial category for observed-location prediction, not a generalizable feature.

Table 3. Validation Results — All Candidates

model	seconds	MAE	RMSE	R2
ridge_trend	0.0137	0.2267	0.2808	0.8356
forest	0.1018	0.2929	0.3571	0.7342
tree	0.0127	0.4032	0.4793	0.5212
median	0.0161	0.6463	0.7881	-0.2947

Selected model: ridge_trend (lowest validation MAE, ties retain the earlier/simpler candidate by stable sort). Ridge with a standardized year term outperforms both the median baseline and the tree/forest — consistent with the underlying generator's additive linear year trend plus fixed per-district/season offsets, exactly the structure ridge is designed to fit.




Figure R1. Validation MAE comparison across median, ridge_trend, tree and forest.

6. Development-Only Rolling-Origin Stability (Version 3.0 Addition)

Table 4. Three One-Year Rolling Origins (ridge_trend only, for brevity — full table in CSV)

origin	model	training_max_year	n	MAE	RMSE	R2
2014	ridge_trend	2013	30	0.2136	0.2655	0.8678
2015	ridge_trend	2014	30	0.2074	0.2637	0.8096
2016	ridge_trend	2015	30	0.2341	0.2859	0.8616

All three origins select ridge_trend as materially the strongest model (MAE 0.207-0.234 t/ha), consistent with the final validation/test selection — this is descriptive stability evidence only, not a confidence interval (manual Section 7).

7. Locked-Test Results (Evaluated Once)

Table 5. Test Results — Selected Model vs. Predeclared Median Baseline

model	MAE	RMSE	R2	batch_inference_seconds
ridge_trend	0.2052	0.2579	0.8667	0.0053

model	MAE	RMSE	R2	batch_inference_seconds
median	0.6365	0.8109	-0.3178	




Figure R2. Locked-test actual vs. predicted yield (t/ha), with the identity line marking perfect prediction.




Figure R3. Locked-test residuals (actual - predicted) vs. predicted yield.

Table 6. Per-Year Robustness (Descriptive, Not a Confidence Interval)

year	n	MAE	RMSE	R2
2017	30	0.2136	0.2575	0.8683
2018	30	0.1968	0.2583	0.8651

MAE is nearly identical across the two locked-test years (0.214 vs 0.197 t/ha), so neither test year dominates the pooled result reported in Table 5.

8. Five Largest Regression Errors

row_id	state	district	season	year	yield_t_ha	prediction	error
R00390	Tamil Nadu	Thanjavur	Rabi	2017	3.1830	2.5310	0.6520
R00321	Uttar Pradesh	Kanpur	Kharif	2018	5.0300	4.4586	0.5714
R00110	West Bengal	Murshidabad	Rabi	2017	3.6290	3.0844	0.5446
R00055	West Bengal	Hooghly	Rabi	2018	4.5180	4.0018	0.5162
R00307	Uttar Pradesh	Lucknow	Rabi	2018	2.7200	3.2143	0.4943

Suspected causes (hypotheses, per Task 8): all five largest errors are in Rabi-season, later-year (2017/2018) observations with the smallest per-district training history (each district contributes only 20 training rows), consistent with the manual's guidance that limited years and district-level aggregation are known limitations — absent irrigation/variety data can suggest, but not confirm, an explanation (Section 7).

9. Advanced Learner Extension 1 — Crop-Label Classification

Training Data Exploration




Figure C0a. Training-set crop-label counts.




Figure C0b. Training-set soil pH distribution.

Class counts are balanced by construction (130 rows/class before splitting), so majority-class imbalance is not a confound in this run's classification metrics — a genuine limitation, since the real D3 Kaggle export's class balance is explicitly unverified per the manual's Section 2 dataset table. Soil pH clusters into overlapping but distinguishable bands per crop (chickpea centred near 7.2, banana/rice nearer 6.0-6.2), which is consistent with pH appearing as a moderately informative feature in the fitted forest.

Table 7. Validation Results — Majority / Logistic / Forest

model	seconds	accuracy	macro_precision	macro_recall	macro_F1
forest	0.1146	0.9725	0.9735	0.9725	0.9726
logistic	0.0108	0.9615	0.9615	0.9615	0.9614
majority	0.0056	0.1429	0.0204	0.1429	0.0357

Selected model: forest (highest macro-F1). The 60-tree random forest separates the seven synthetic crop profiles almost perfectly, well ahead of the majority baseline (macro-F1 0.036) and modestly ahead of logistic regression.




Figure C1. Validation macro-F1 comparison across majority, logistic and forest.

Table 8. Locked-Test Results — Selected Model vs. Predeclared Majority Baseline

model	accuracy	macro_precision	macro_recall	macro_F1	batch_inference_seconds
forest	0.9560	0.9566	0.9560	0.9561	0.0152
majority	0.1429	0.0204	0.1429	0.0357	




Figure C2. Locked-test confusion matrix across the seven crop labels.

Table 9. Per-Class Precision / Recall / F1 (Locked Test)

Class	Precision	Recall	F1	Support
banana	0.963	1.000	0.981	26
chickpea	1.000	1.000	1.000	26
coffee	0.885	0.885	0.885	26
cotton	0.960	0.923	0.941	26
maize	0.889	0.923	0.906	26
mango	1.000	1.000	1.000	26
rice	1.000	0.962	0.980	26

The weakest classes are coffee (F1 0.885) and maize (F1 0.906), which share overlapping temperature/humidity/rainfall ranges in the synthetic profiles (Figure C2 shows most confusion is between these two); chickpea and mango are separated perfectly (F1 1.0) due to their distinct pH and rainfall bands. Macro F1 equally weights all seven classes, so this correctly surfaces the coffee/maize weakness that plain accuracy would understate.

Table 10. Five Locked-Test Misclassifications

row_id	N	P	K	temperature	humidity	pH	rainfall	label	prediction	error
C00420	104.9900	26.4100	38.8700	24.3200	62.9100	6.5800	64.0600	cotton	coffee	1
C00598	78.8000	33.1300	18.0100	17.7100	47.5600	6.4100	133.8100	coffee	maize	1
C00431	124.8300	69.7700	61.3500	22.2600	84.9500	7.2700	54.0900	cotton	banana	1
C00631	84.3300	30.8700	21.0100	23.3300	79.5300	6.6400	113.4400	coffee	maize	1
C00142	91.4900	23.3100	31.5100	23.5700	62.3200	7.4600	127.9100	maize	coffee	1

10. Advanced Learner Extension 2 — Depth Ablation

Adding a shallower forest_depth6 candidate alongside the core models, evaluated on the identical training/validation split (advanced=true, Appendix B):

model	seconds	MAE	RMSE	R2
ridge_trend	0.0100	0.2267	0.2808	0.8356
forest	0.0855	0.2929	0.3571	0.7342
forest_depth6	0.0839	0.3336	0.4075	0.6539
tree	0.0093	0.4032	0.4793	0.5212
median	0.0110	0.6463	0.7881	-0.2947

The depth-6 forest is worse than the depth-12 forest (MAE 0.334 vs 0.293) and both remain worse than ridge_trend — reducing forest depth does not close the gap to a model that already matches the data's true additive/linear structure; it only reduces the forest's capacity to (partially) approximate that structure through splits.

11. Reproducibility and Acceptance

Check	Regression	Classification
Split disjoint (chronological/stratified)	true	true
Reload consistent (bundle predictions match in-memory)	true	true
Valid predictions (finite / within known label set)	true	true
Data hash matched at test time	true	true
Failure rows available for audit	5	5

Environment: Python 3.12.3, scikit-learn 1.8.0, numpy 2.4.4, pandas 3.0.2, seed=42. The TEST_LOCK file is written before either locked-test evaluation and the code re-verifies the data/config SHA-256 hashes and exact split membership before scoring, so the test set cannot have been altered or re-selected after being opened.

12. Discussion Questions — Answers

1. What changes when the prediction unit moves from district to farm?

Farm-level prediction would need farm-specific identifiers, far more granular and likely noisier yield records, and features like plot size, soil test results and irrigation source that are meaningless or unavailable at district aggregation — the district-level model here cannot be assumed to generalize down to individual farms.

2. Why does a production predictor invalidate a yield-derived target?

Because $\text{yield_t_ha} = \text{production_t} / \text{area_ha}$, including production as an input lets the model algebraically reconstruct most of the target rather than learn a genuine relationship, producing an unrealistically low error that would not hold if production were unknown at prediction time.

3. At what date must rainfall be available for a pre-sowing model?

Before the sowing decision for that season/year — using whole-season rainfall (which is only fully known after harvest) would leak future information, which is why the core regression whitelist excludes rainfall entirely (Section 3, canonical field table).

4. Why might a forest fail to beat the median?

If the true relationship is close to a simple additive trend plus fixed offsets, a forest's flexible partitioning adds variance without capturing structure a simpler model already captures well — though in this run the forest still clearly beats the median (MAE 0.293 vs 0.646), it does not beat ridge for exactly this reason.

5. What does a negative R^2 tell you about this evaluation set?

That the model performs worse than simply predicting the mean of that evaluation set — seen here for the median baseline on the locked test ($R^2 = -0.318$), which is expected since it ignores the district/year trend other models use.

6. How can a source revision invalidate a historical backtest claim?

If official crop statistics are later revised (common for agricultural series), a backtest using the current, revised numbers implicitly uses information not available at the historical prediction time, inflating apparent accuracy.

7. Why does chronological splitting not prove spatial generalization?

A chronological split only tests whether the model generalizes forward in time for the same set of districts it was trained on — it says nothing about districts absent from training, which is why the manual requires a separate spatial-shift extension (Section 9) to test that.

8. What information would justify keeping two repeated keys?

Documented, legitimate re-reporting or aggregation (e.g., a corrected revision explicitly superseding an earlier entry) — not silent duplication; the pipeline instead rejects any repeated district-season-year key outright and requires source review (Section 6, Task 2).

9. How would district boundary changes alter grouping?

A district split or renamed between years would fragment or merge historical series under different names, making year-over-year district comparisons invalid unless boundaries are harmonized first — a known D1/D2 limitation (Section 2).

10. Which agricultural consequence makes MAE alone insufficient?

MAE treats a 0.5 t/ha error the same everywhere, but the same absolute error may be agronomically trivial for a high-yield district and severe for a low-yield, food-insecure one — economic/agronomic thresholds require domain input the manual explicitly does not impose (Section 1).

11. Why are two test years insufficient for a strong temporal CI?

A confidence interval needs enough independent-ish samples to estimate variability; two year-batches give only two correlated data points for between-year variation, nowhere near enough for a defensible interval (Section 7).

12. When is a classification probability unsuitable as confidence?

When the model has not been calibrated (i.e., predicted probabilities have not been checked against observed frequencies) — the manual explicitly warns against presenting maximum class probability as validated confidence without calibration (Section 9).

13. How can apparently balanced augmented data still be biased?

If augmentation duplicates or perturbs a small set of original records to fill out class counts, per-class support can look balanced while the underlying diversity and label-generation provenance remain narrow or unverified — a documented D3 limitation (Section 2).

14. Which additional measurements are necessary for agronomic crop selection?

Soil testing, irrigation availability, crop variety, economics, weather forecasts and local expert evidence — none of which are in this exercise's scope, which is why the manual insists crop-label output is a dataset label, not a farming prescription (Section 9, 13).

15. What evidence would justify deployment beyond this exercise?

A verified real-data run with instructor-checked provenance, a timed classroom/field pilot, calibrated uncertainty intervals, spatial-generalization testing, and expert agronomic review — all explicitly listed as not yet completed in the manual's Section 16 verification summary.

13. Viva Questions and Key Points (Full 25-Question Set)

Question	Key point
What is the target?	Rice yield in verified t/ha for district-season-year.
Regression or classification?	Yield is continuous; crop label is categorical.
Why one crop?	Consistent units and more coherent prediction scope.
Why median baseline?	Optimal constant for training absolute-error loss.
How does the tree predict?	Split feature space, then use training leaf means.
What does the forest average?	Predictions from randomized bootstrap-trained trees.
Why limit depth?	Bound complexity and reduce overfitting.
Why exclude production?	With area it reconstructs the target.
Can area always be used?	Only if available at prediction time; core excludes it.
What is training-only fitting?	Learn transformations on training rows, apply elsewhere.
Why a pipeline?	Keeps preprocessing tied to the fitted model, fit only on training rows.
What is SST?	Sum of squared deviations from the evaluation-target mean.
RMSE versus MAE?	RMSE penalizes large errors more strongly.
Can R^2 be negative?	Yes — worse squared error than the evaluation-set mean reference (seen for the median baseline here).
Why omit MAPE?	Zero/small yields cause undefined or unstable ratios.
Why reserve latest years?	Simulate later-year generalization.
What does validation choose?	Model/settings without using the final test — selection here used validation MAE only.
Is year-feature regression a time-series comparison?	No; temporal data alone do not establish course outcome CO6.
Why save row IDs?	Audit exact split membership and overlap between runs.
What is macro F1?	Equal-weight mean of per-class F1 values — used because accuracy alone would hide the coffee/maize confusion.
What is calibration?	Agreement of predicted probabilities and observed frequencies.
What does a residual show?	Actual minus predicted; sign gives bias direction — positive residuals here mean underprediction.
Why reject unknown districts?	Their reliability was not established by this task; predict_validated() enforces this.
Does importance imply causality?	No — it describes fitted predictive association only.
What proves reproducibility?	Same data hash, environment, config, splits and reload outputs — all verified in Section 11.

14. Common Mistakes — Self-Audit

Common mistake	How this submission addresses it
Use production as input	Avoided — REG whitelist is exactly [state, district, season, year]; production/area never enter X.
Randomly split future-year task	Avoided — ordered year blocks enforced and asserted in code (Section 3).
Impute before splitting	Avoided — imputer/scaler/encoder live inside the sklearn Pipeline, fit only on training rows.
Tune on test residuals	Avoided — TEST_LOCK is written before scoring; selection.json is fixed at validation time and hash-verified unchanged at test time.

Common mistake	How this submission addresses it
Report only accuracy for classification	Avoided — macro precision/recall/F1 and per-class metrics all reported (Table 7-9), surfacing the coffee/maize weakness accuracy alone would hide.
Claim forest extrapolation	Avoided — Section 6 discussion explicitly notes tree ensembles do not extrapolate the year trend; this is exactly why ridge wins here.
Call feature importance causal	Avoided — Section 12, discussion question 15 explicitly frames importance as association only.
Load unknown serialized model	Avoided — only the locally-saved, self-produced <code>selected_bundle.joblib</code> is reloaded in Task 8.

15. Responsible Agricultural Analytics

- No farm coordinates, farmer identifiers or other personal data are present; district/state are aggregate administrative categories only.
- Yield predictions are estimates for a student-analyst regional-planning exercise, not proof of future yield or an agronomic recommendation.
- Crop-label classification output is a dataset label reflecting the synthetic generator's profiles, not farming advice — a real deployment would need soil testing, irrigation, variety, economic and expert agronomic input per Section 13.
- `predict_validated()` rejects unknown districts/categories and out-of-range years by design, so the demonstration inference function will not silently extrapolate to unseen locations or future years.
- Only the locally-produced `selected_bundle.joblib` is loaded in this notebook; the manual's warning against loading untrusted serialized models is followed throughout.

Risk Register (Manual Section 13)

Risk / cause	Impact	Mitigation	Responsible stakeholder
Unit mismatch / undocumented conversion	Misleading yield estimates	Source-unit approval and conversion tests	Instructor/data steward
Temporal leakage / revised statistics	Inflated backtest performance	Availability audit; qualify vintage limitation	Student analyst
Sparse regions / biased coverage	Poor regional decisions	Support counts and regional validation	Analyst/agronomist
Label misuse / unclear provenance	Unsuitable crop advice	Educational labeling and expert validation	Instructor/demo owner
Weather or policy shift / drift	Degrading predictions	Monitor later-year errors and revalidate	Future system owner
Untrusted model loading	Code execution risk	Load only trusted saved artifacts	Lab administrator

16. Assessment Rubric Self-Check (100 Marks)

Criterion	Marks	Evidence in this submission
Framing	8	Prediction contract with unit, horizon and excluded fields (Section 1).
Source and unit audit	10	Dataset governance record with conversion/exclusion notes (Section 2).
EDA	8	Training-only target/year distribution figures (Section 5, Figure R1; classification Figure C1).
Preprocessing	10	OneHotEncoder + median-impute/StandardScaler inside a training-only-fitted Pipeline (Section 4).
Leakage and split	14	Ordered year blocks with code-level assertions; stratified classification split with independence note (Section 3).
Baseline	8	Median/majority baselines implemented and reported on both validation and locked test (Tables 3, 5, 7, 8).
Models and comparison	10	Ridge/tree/forest and logistic/forest compared on identical splits (Tables 3, 7).
Selection and test	10	Validation-only selection, one-time TEST_LOCK, hash-verified before scoring (Section 7, Section 11).
Error and robustness	10	Five-case audits for both tasks plus per-year robustness and rolling origins (Sections 6, 8, 10).
Responsible interpretation	5	Risk register and scope/limitation statements (Section 15).

Criterion	Mark s	Evidence in this submission
Reproducibility/report	7	Reload checks, saved versions, this report (Section 11).

17. Evidence-Based Benchmarking Note

The manual itself is an internal design review (Version 3.0, self-scored 8.21/10 against a Version 2.0 baseline of 7.50/10 across 12 critique dimensions), not an independently validated instrument. This submission does not add new external validation of the manual; it only reports where the executed pipeline stands against the manual's own stated critique dimensions, using this run's actual evidence rather than assumed compliance.

Dimension	Status for this submission
Dataset readiness	Real file bytes for D1/D2/D3/D5 remain inaccessible in this environment; synthetic schema-matched substitutes used throughout (Section 1).
Leakage prevention	Ordered year blocks, training-only pipeline fitting, and encoder-category assertions all verified true at runtime (Section 11).
Model selection	Ridge/tree/forest and logistic/forest both compared under identical splits; selection uses validation-only metric with no test peeking (Sections 5, 9).
Statistical rigor	Three development rolling origins reported (Section 6) as descriptive stability evidence; the manual is explicit this does not establish a confidence interval — no such claim is made here.
Reproducibility	Full acceptance/version/hash record captured (Section 11); this run itself constitutes the "real notebook run" the manual's Section 16 flags as outstanding, on synthetic rather than real data.

18. Advanced Uncertainty Design (Manual Section 7)

The manual specifies that a credible advanced uncertainty design needs at least five development-year rolling origins (this submission reports three, the Version 3.0 core minimum — Section 6, Table 4) and that any block bootstrap requires an explicit block length with enough independent blocks; an ordinary row-level bootstrap must not be labeled valid under regional/time dependence. This submission does not attempt a bootstrap confidence interval for exactly this reason: with only two locked-test years (Table 6), any row-level resampling would treat clearly dependent same-year observations as independent, producing an artificially narrow and misleading interval. Extending the rolling-origin count from three to five or more development-year cutoffs, and only then considering a properly block-structured bootstrap, is listed as future work rather than attempted here.

19. References

- Government of India, District-wise season-wise crop production statistics catalog (D1).
- TCI/Cornell, District-Level Database for Indian Agriculture and Allied Sectors (D2).
- ICRISAT, Crops dataset description (D2/D5).
- Souryabrata Mohapatra, ICRISAT District-Level Data: Heterogeneous Climate Effect on Crop Yield and Associated Risks to Water Security in India, DOI 10.17632/ywp3y5j9vv.1 (D5).
- Atharva Ingle, Crop Recommendation Dataset, Kaggle (D3).
- FAO, FAOSTAT, Crops and livestock products (D4).
- Stanford University, CS229 Final Report Guidelines.
- MIT, 6.390 Spring 2023 Grading, Lateness, and Extensions.
- scikit-learn, Common pitfalls and recommended practices.

20. Conclusion

This submission implements the complete Experiment 08 three-hour core plus both listed Appendix-A-supported extensions: crop-label classification and depth ablation. The regression core selected ridge_trend by validation MAE, confirmed it on a single locked-test pass (MAE 0.205 t/ha, $R^2=0.867$) against the predeclared median baseline, and corroborated the choice with three development-only rolling origins and near-identical per-year test robustness. The classification extension selected a 60-tree random forest (test macro-F1 0.956), correctly surfacing per-class weaknesses (coffee, maize) that plain accuracy

would have hidden. The depth-ablation extension confirmed that reducing forest depth does not close the gap to ridge on this data. All results, artifacts and figures were produced by running the manual's unmodified lab08.py, and every leakage/reproducibility assertion embedded in that code (ordered splits, training-only fitting, one-time test lock, hash verification, reload consistency) passed without modification.

As with the regression/classification data-source caveat in Section 1, the single largest limitation is that these are synthetic, schema-matched substitutes for the manual's preferred real sources (D5 for rice, D3 for crop labels); re-running this identical notebook against instructor-verified real exports is the necessary next step before any of these numbers could be reported as genuine agricultural findings.